

ML Part 2 tutorial

Dimensionality reduction & cross-validation

Jérôme Dockès & Nikhil Bhagwat

QLS course 2021-07-30



Problem setting

$$Y = f(X) + E \quad (1)$$

- $Y \in \mathbb{R}$: output (a.k.a. target, dependent variable) to predict
- $X \in \mathbb{R}^p$: features (a.k.a. inputs, regressors, descriptors, independent variables)
- $E \in \mathbb{R}$: unmodelled noise
- f : the function we try to approximate

Parameter estimation a.k.a. model fitting

Minimize a sum of:

- the empirical risk: error on training data
- a regularization term

Example: logistic regression

$$\operatorname{argmin}_{\beta, \beta_0} \frac{1}{2} \|\beta\|_2^2 + C \sum_{i=1}^n \log(\exp(-y_i (\mathbf{X}_i^T \beta + \beta_0)) + 1) \quad (2)$$

- β, β_0 : parameters to be *estimated*
- C : hyperparameter, *chosen* prior to learning (controls amount of regularization)

`sklearn.linear_model.LogisticRegression`

scikit-learn "estimator API": fit; predict

```
estimator = LogisticRegression(C=1)
estimator.fit(X_train, y_train)
predictions = estimator.predict(X_test)
```

https://scikit-learn.org/stable/getting_started.html
`sklearn.linear_model.LogisticRegression`

Evaluating performance with `sklearn.metrics`











```
estimator = LogisticRegression(C=1)
estimator.fit(X_train, y_train)
predictions = estimator.predict(X_test)

accuracy = metrics.accuracy_score(y_test, predictions)
```

https://scikit-learn.org/stable/getting_started.html
`sklearn.linear_model.LogisticRegression`
`sklearn.metrics` more info on model evaluation

`ex_01_fit_predict.py`

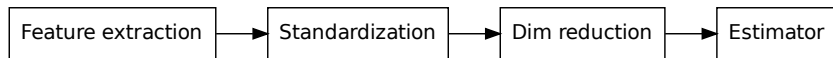
Cross-validation

Fold 0	Train		Score 0
	Test		
Fold 1	Train		Score 1
	Test		
Fold 2	Train		Score 2
	Test		
Fold 3	Train		Score 3
	Test		
Fold 4	Train		Score 4
	Test		

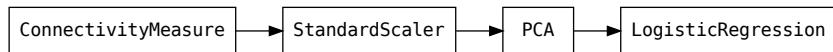
scikit-learn.org/stable/modules/cross_validation.html
`sklearn.model_selection.cross_validate`
`ex_02_cross_validate.py`

Dataset transformations

Typical pipeline



Example: for autism prediction with fMRI from ML part 1



scikit-learn "transformer API": `fit`; `transform`

```
transformer = StandardScaler()  
transformer.fit(X_train)  
transformed_X = transformer.transform(X_train)
```

can also be written:


```
transformer = StandardScaler()  
transformed_X = transformer.fit_transform(X_train)
```

[sklearn.preprocessing.StandardScaler](#)
[scikit-learn "getting started"](#)
[scikit-learn "user guide"](#)

`ex_03_transformer.py`

scikit-learn "transformer API": fit; transform

```
transformer = StandardScaler()  
transformed_X = transformer.fit_transform(X_train)  
  
transformed_X_test = transformer.transform(X_test)
```



[sklearn.preprocessing.StandardScaler](#)
[scikit-learn "getting started"](#)
[scikit-learn "user guide"](#)

Example: `preprocessing.StandardScaler`

`fit:`

Compute mean and standard deviation of each column

`transform:`

Subtract mean and divide by standard deviation

`sklearn.preprocessing.StandardScaler`

Example: `feature_selection.SelectKBest`

`fit:`

- compute ANOVA or correlation for each column of X
- Remember the indices of the k columns with highest scores

`transform:`

- Index input to keep only the k selected columns

`sklearn.feature_selection.SelectKBest`

https://scikit-learn.org/stable/modules/feature_selection.html

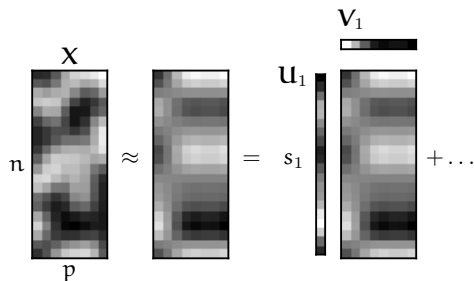
`ex_04_feature_selection.py`

Example: decomposition.PCA

fit:

- Compute Singular Value Decomposition of X

$$X = U S V^T \quad (3)$$



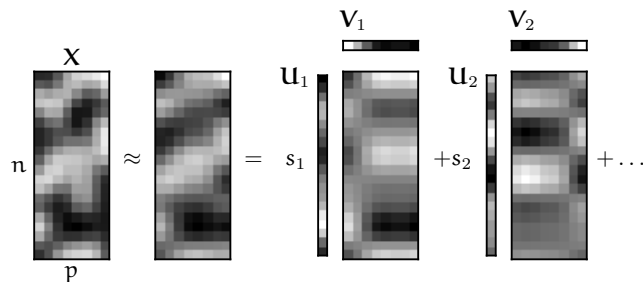
Explained variance: 0.53

Example: decomposition.PCA

fit:

- Compute Singular Value Decomposition of X

$$X = U S V^T \quad (4)$$




Explained variance: 0.84

Example: decomposition.PCA

fit:

- Compute Singular Value Decomposition of X

$$X = U S V^T \quad (5)$$



Explained variance: 0.97

Example: `decomposition.PCA`

`fit:`

- Compute Singular Value Decomposition of X

$$X = U S V^T$$

- store V

`transform:`

Compute projection on column space of V : simply multiply by V^T

Notes

- `fit_transform`: simply return $U S$
- V^T is the `'components_'` attribute of a fitted `'PCA'` instance

`sklearn.decomposition.PCA`





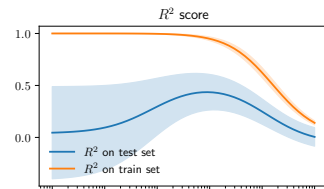
Chaining transformations

Use `sklearn.pipeline.Pipeline` or
`sklearn.pipeline.make_pipeline`:

```
pipe = make_pipeline(  
    standardizer, dim_reductor, estimator  
)  
pipe.fit(X, y)
```

Example:









```
make_pipeline(  
    StandardScaling(), PCA(), LogisticRegression()  
)
```












$$\text{Var}(\hat{\beta}_i) = \mathbb{E}(\hat{\beta}_i - \mathbb{E}(\hat{\beta}_i))^2$$









$$\text{Bias}(\hat{\beta}_i) = \mathbb{E}(\hat{\beta}_i) - \beta_i$$

Nested cross-validation






Fold 0	Train	Fold 0	Train	For all λ	
			Test	For all λ	
		Fold 1	Train	For all λ	
			Test	For all λ	
		Fold 2	Train	For all λ	
			Test	For all λ	
		Refit		For best λ	
	Test				
					Score 0

Fold 1	Train	Fold 0	Train	For all λ	
			Test	For all λ	
		Fold 1	Train	For all λ	
			Test	For all λ	
		Fold 2	Train	For all λ	
			Test	For all λ	
		Refit		For best λ	
	Test				

Score 1

Fold 2	Train	Fold 0	Train	For all λ	
			Test	For all λ	
		Fold 1	Train	For all λ	
			Test	For all λ	
		Fold 2	Train	For all λ	
			Test	For all λ	
		Refit		For best λ	
	Test				

Score 2

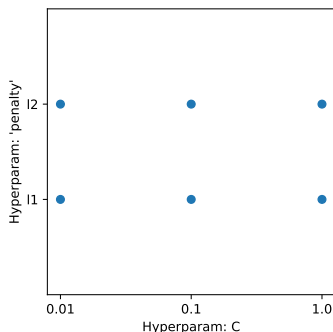
Fold 3	Train	Fold 0	Train	For all λ	
			Test	For all λ	
		Fold 1	Train	For all λ	
			Test	For all λ	
		Fold 2	Train	For all λ	
			Test	For all λ	

Implementing nested CV

`ex_05_nested_cross_validation.py`

Cross-validation and hyperparameter selection in scikit-learn

- `sklearn.pipeline.Pipeline` or `sklearn.pipeline.make_pipeline`
- `sklearn.model_selection.GridSearchCV`
- `sklearn.model_selection.cross_validate`
- use *CV estimators! `RidgeCV`, `LogisticRegressionCV`, ...



Cross-validation pitfalls

- fitting part of the pipeline on the whole data: use Pipeline
- ignoring some dependencies in the data: use the appropriate cv iterator:
https://scikit-learn.org/stable/modules/cross_validation.html#cross-validation-iterators
- good cv scores on one dataset do not guarantee generalization to new data