

# Machine learning Part 2

Jérôme Dockès

QLS course 2021-07-30



**McGill**  
UNIVERSITY



# Recap of part 1

## Supervised learning

- Regression: least-squares linear regression
- Classification: logistic regression

## Regularization

- $\ell_2$  a.k.a. ridge regularization

## Model evaluation and selection

- Out-of-sample generalization; independent test set
- Performance metrics:
  - regression: mean squared error
  - classification: accuracy, ROC curve
- Cross-validation

Don't remember? watch Part 1 again!

# Notation & vocabulary

## Supervised learning framework

$$Y = f(X) + E \quad (1)$$

- $Y \in \mathbb{R}$ : output (a.k.a. target, dependent variable) to predict
- $X \in \mathbb{R}^p$ : features (a.k.a. inputs, regressors, descriptors, independent variables)
- $E \in \mathbb{R}$ : unmodelled noise
- $f$ : the function we try to approximate

## Example: linear regression

$$Y = \beta_0 + \langle X, \beta \rangle + E \quad (2)$$

$$= \beta_0 + \sum_{j=1}^p X_j \beta_j + E \quad (3)$$

"learning" = estimating  $\beta_0 \in \mathbb{R}$  and  $\beta \in \mathbb{R}^p$

# Dimensionality reduction

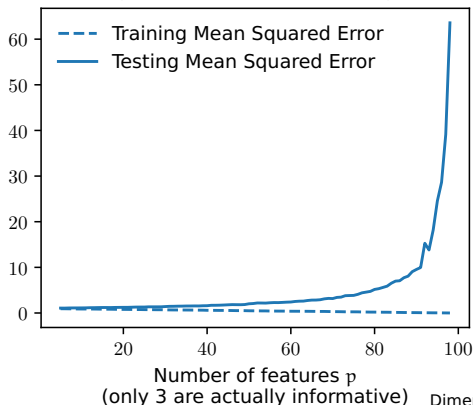
Problems when the number of features  $p$  becomes large

- Bigger errors on test data (larger variance of predictions)
- Numerical stability issues
- Computational cost and memory usage

# Model complexity: overfitting

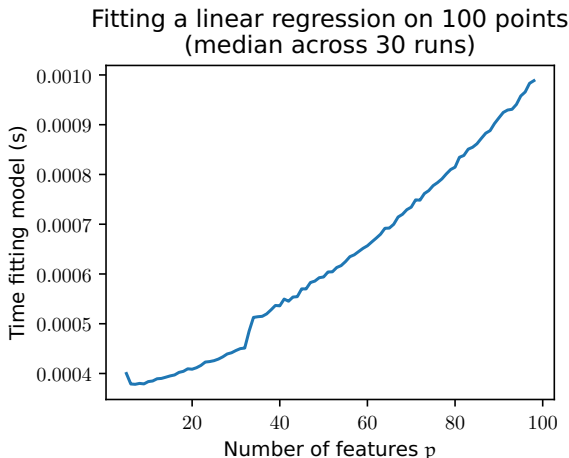
- Model complexity increases with dimension.
- Example: a linear model in dimension  $p$  can fit exactly (0 training error) any set of  $p + 1$  points.
- Risk of overfitting: fitting exactly training data but failing on test data

Fitting a linear regression on 100 points  
(median across 30 runs)



# Cost of fitting many parameters

- Many algorithms require polynomial time in  $p$
- Implementations often make copies of the design matrix (e.g. for centering & rescaling)



# Univariate feature selection

- a.k.a. feature screening, filtering ...
- Check features (columns of  $X$ ) one by one for association with the output  $y$
- Keep only a fixed number or percentage of the features

## Simple (linear) association criteria

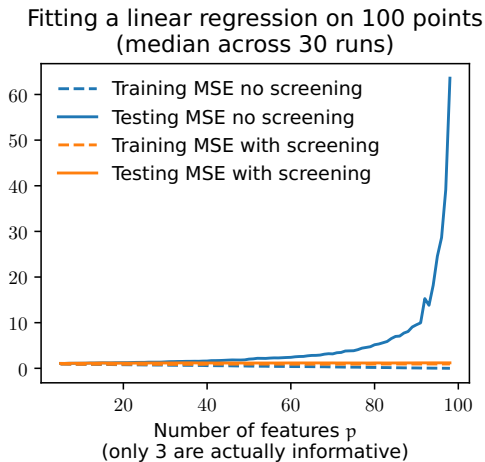
- for regression: correlation
- for classification: ANalysis Of VAriance

Read more in the scikit-learn user guide

[https://scikit-learn.org/stable/modules/feature\\_selection.html#feature-selection](https://scikit-learn.org/stable/modules/feature_selection.html#feature-selection)

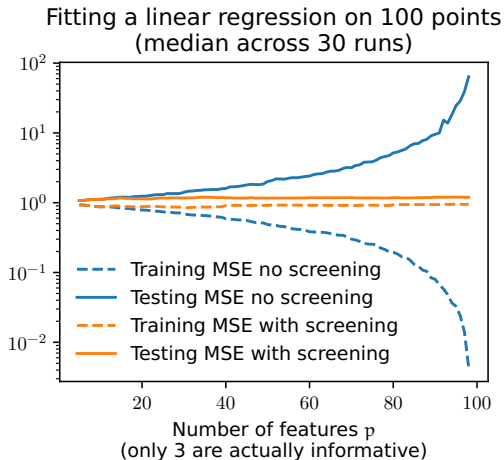
# Univariate feature selection

Keeping only the 10 best features (most correlated with  $y$ )



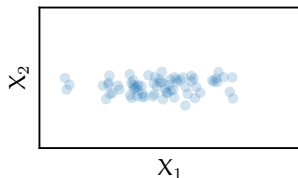


# Same plot in log scale

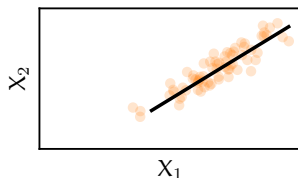


# Linear decomposition methods

Maybe OK to drop  $X_2$ :



Data low-dimensional but no feature can be dropped:



Find a better referential in which to represent the data

# Linear regression: projection on the column space of $X$

Approximate  $y$  as a combination of the columns of  $X$

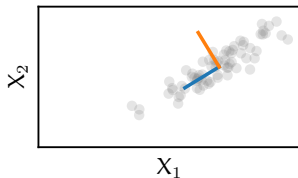
$$\hat{y} = X \hat{\beta} \in \mathbb{R}^n \quad (4)$$

- The columns of  $X$  are a family of  $p$   $n$ -dimensional vectors
- When  $p$  is high or the columns of  $X$  are correlated, we want to use a family of  $k < p$  instead
- Feature selection: drop some columns, keep only  $k$
- Could we build a better family of  $k$  vectors?

# Principal Components Regression

- Approximation of  $X$  of rank  $k$ : find a family of  $k$  basis vectors and approximate each column of  $X$  as a mixture of these  $k$  vectors
- Find the family that gives the best approximation: the one with the smallest Frobenius norm of the reconstruction error.
- This is the same as finding the  $k$  orthogonal directions in which  $X$  varies the most

# Principal Components: feature space



# Ridge regression and PCA

- Both ridge regression and PC regression compute the coordinates of  $y$  in the basis given by the SVD of  $X$
- ridge shrinks the coefficients of sv  $d_j$  by a factor  $d_j^2 / (d_j^2 + \lambda)$
- PC regression sets the coefficient to 0 for all but the  $k$  largest  $d_j$

# Other decomposition methods

- Take  $y$  into account
- Different criteria (sparsity, independence, ...)

# Nested cross-validation: setting hyperparameters

How can we choose:

- Number of features or PCA components  $k$ ?
- The ridge hyperparameter  $\lambda$ ?

Try a few and pick the best one... But measure its performance on separate data!



# Some common pitfalls with cross-validation

- Ignoring dependencies between samples
- Ignoring dependencies between CV scores
- Over-interpreting good CV scores

Two sources of variance: training data and test sample

Don't use Leave-One-Out Cross-validation

# fMRI decoding

- Describe data and task

# The decoding pipeline

- Masking: extracting voxels that are inside the brain
- Feature selection with ANOVA
- Classifier: logistic regression

# Implementation: in class